

HIRESHEMA

AI Recruiting Agent

for Job Seekers

Daily AI-curated job matches scored against your resume and career paths — delivered to your dashboard every morning.

Business review · Tech stack · Pricing & unit economics

Prepared for: Founding partner

Why we are having this meeting

THE PRODUCT

Hireschema is a daily job matching agent. Users upload a resume, set 2-5 career paths, and every morning we deliver a small, ranked set of fresh roles already scored for fit — plus AI tools to apply faster.

THE USER

Active and passive job seekers — operations, customer success, product, and engineering professionals who already spend 30-60 min/day filtering low-fit listings on LinkedIn and Indeed.

WHAT THIS MEETING DECIDES

- Pricing tier — lock the Pro price point and validate the Free !' Pro funnel.
- Unit economics — confirm the cost per Pro user and what margin we are protecting.
- Break-even — agree on the paying-user count where we cover fixed costs.

Job search is a daily inbox of irrelevance

- 40+ min daily filter time — average active job seeker scrolls 100+ listings before finding 1-2 worth applying to.
- 70% off-target listings — existing platforms optimize for advertiser revenue, not match quality.
- 1-in-50 reply rate — generic applications hit single-digit response rates. The pain is finding the right jobs, fast.

The job is to compress the daily filter into a 2-minute inbox. Curation, not catalog.

A daily curated inbox of the 10 best-fit roles

Hireschema runs a personal recruiter agent for every user, every morning. We pull live listings from career-site APIs (no ads, no resume harvesting), rank them with AI against your resume + career paths, and deliver only the top 10 to your dashboard.

Once you find a fit, the same agent generates the cover letter, tailors your resume, and preps you for the interview — without you opening another tool.

Core capabilities

- Daily 10 matches — ranked 0–100 against your profile
- AI cover letter — one click, role-specific, anti-slop
- Resume tailoring — auto-emphasizes relevant experience
- Interview prep — tough questions per role, with model answers
- Application tracker — Save / Applied / Interviewing pipeline

What the user sees, end to end

Concern	Tech	Why this choice
Framework	React 19 + TypeScript	Strict mode, hooks-only architecture, no class components.
Build	Vite 6	ESM-first dev server, sub-second HMR, ~8s production build.
Styling	TailwindCSS 4 + Shadcn UI	Utility-first, accessible primitives (Radix).
Routing	React Router v7	File-based routes, Suspense-friendly, server-data ready.
State	React Context + Firestore onSnapshot	Real-time profile + jobs sync, no Redux/Zustand.
Auth	Firebase Auth (Google OAuth)	ID tokens verified server-side via Firebase Admin.
Animations	Framer Motion	Modal + panel transitions, scroll-driven reveals.
Icons	Lucide React	Tree-shakeable SVG icon set.
Toasts	Sonner	Lightweight, accessible notifications.
Markdown	react-markdown + remark-gfm	Renders AI-generated content safely.

Pipeline, storage, security, observability

Concern	Tech	Notes
Runtime	Vercel Serverless Functions (Node 20, TypeScript)	On-demand only — zero idle cost.
Long jobs	GitHub Actions runners (ubuntu-latest)	Workflow runs the multi-minute job pipeline outside Vercel's 60s cap.
Auth	Firebase Admin SDK	verifyIdToken() on every protected endpoint.
DB	Firebase Firestore (named DB)	users/, daily_matches/, cronRuns/, trackedJobs/, blog_posts/.
Discovery	Apify HTTP API — Career Site Actor (fantastic-jobs)	Strict + broad pulls, deduped by fingerprint.
AI	OpenRouter via openai npm SDK	Single proxy, swap models without code change.
Cron	GitHub Actions schedule (8 AM + 1 PM IST) + dashboard auto-trigger	Three independent reliability layers.
Security	Firestore rules + INTERNAL_CRON_SECRET + CRON_SECRET	Admin-SDK writes bypass rules; client reads only own docs.
Observability	Console logs to Vercel + GitHub Actions, cronRuns Firestore audit trail	Per-user run status + failureReason.

Every model we call, what it does, what it costs

Task	Model	Purpose	Cost
Job scoring	openai/gpt-4o-mini	Scores top 30 candidates 0-100 against resume + career paths.	\$0.15 / 1M in · \$0.60 / 1M out
Career path suggestion	openai/gpt-4o-mini	Reads resume + structured profile, proposes 5-7 target roles.	\$0.15 / 1M in · \$0.60 / 1M out
Resume analysis	openai/gpt-4o-mini	Extracts skills, seniority, salary band, location flexibility.	\$0.15 / 1M in · \$0.60 / 1M out
Cover letter	openai/gpt-4o	120-word personalized letter grounded in resume + role.	\$2.50 / 1M in · \$10 / 1M out
Resume tailoring	openai/gpt-4o	Reorders / surfaces relevant bullets per role. Never invents.	\$2.50 / 1M in · \$10 / 1M out
Interview prep	openai/gpt-4o	10-15 likely questions with model answers per role.	\$2.50 / 1M in · \$10 / 1M out
Job discovery	Apify Career Site Actor	Live listings from 20K+ employer career sites, ATS-aware.	Per actor compute unit

All models are accessed through the code.

Daily delivery is guarded by redundancy, not faith

Pipeline: Discover !' Score !' Curate !' Deliver

- Discover — Apify Career Site Actor (strict + broad pulls, deduped)
- Score — OpenRouter gpt-4o-mini, 85% AI weight + 15% deterministic
- Curate — per-company dedup, top 10 by final score
- Deliver — Firestore write, React dashboard render

Three independent reliability layers

- Primary cron — GitHub Actions runs daily at 8:00 AM IST. Iterates active users, processes each through the pipeline.
- Catch-up cron — Second GH Actions run at 1:00 PM IST. Re-tries any user whose primary run failed or never created a record.
- Auto-trigger — When a user opens the dashboard and finds nothing for today, we dispatch a fresh run automatically.

Big problem, broad audience, willing to pay

- TAM — ~5M active US tech / knowledge-worker job seekers in any given quarter.
- SAM — ~2M English-speaking, remote-friendly job seekers with a digital resume.
- SOM — ~100K free signups, ~5-8% paid conversion = 5,000–8,000 Pro users (3-year horizon).

Willingness-to-pay benchmarks: Teal \$9.99/mo · Jobscan \$49.95/mo · Sonara \$50/mo · Massive \$50/mo · LazyApply \$99/yr.

Where we sit on price and feature depth

	Hireschema	Teal	Jobscan	LazyApply	Sonara
Daily AI matches	' 10 / day	Limited	—	Auto-apply	Auto-apply
AI cover letter	'	'	'	Templated	Templated
Resume tailoring	'	Manual	' ATS	—	—
Interview prep	' AI	—	—	—	—
Monthly price	\$14 (target)	\$9.99	\$49.95	\$8.25 / yr	\$50
Persona	Active seeker	Career planner	ATS-anxious	Mass-applier	Mass-applier

Our wedge: only product th
consumer pricing. Mass-ap

Cost per Pro user per month — small-scale baseline

Service	What we use	Pro / mo	Free / mo	Notes
Apify	Career Site Actor — ~60 runs/user/mo	\$1.50 – \$3.00	\$0.05 – \$0.10	Estimate — confirm with billing
OpenRouter	gpt-4o-mini scoring + gpt-4o tools	\$0.50 – \$1.50	\$0.05 – \$0.15	Token-based; per-user budget
Firebase	Auth + Firestore reads/writes	< \$0.05	< \$0.05	Free tier: 50K reads / 20K writes /
Vercel	Hobby tier — 100 GB bandwidth	\$0	\$0	Covers ~5,000 active users
GitHub Actions	2 cron runs/day per user (~3 min each)	\$0	\$0	Free 2,000 min/mo caps at ~250 Pro
TOTAL	All-in cost basis	\$2.05 – \$4.55	\$0.15 – \$0.35	Estimate — confirm with billing

Free, Pro \$14/mo, Pro Annual \$99/yr

Free — \$0

Acquisition funnel.

- 1 daily match
- Match scoring (basic)
- No AI tools
- Job tracker (basic)

Pro — \$14 / month (RECOMMENDED)

- 10 daily matches
- AI cover letter (unlimited)
- AI resume tailoring
- AI interview prep
- Full job tracker

Pro Annual — \$99 / year (save 41%)

- Everything in Pro
- Priority email support
- Early access to new tools
- Locks in current pricing

Why \$14: 3-4× margin over cost · undercuts Teal feature-for-feature · consumer-friendly (\$9-15 sweet spot).

Revenue, cost, profit at each user-count milestone

Stage	Pro users	Free users	MRR	Cost / mo	Profit / mo	Cum. profit
Pre-launch / friends	10	50	\$140	\$80	\$60	—
Soft launch (Mo 1)	25	200	\$350	\$140	\$210	—
Public launch (Mo 2)	50	500	\$700	\$280	\$420	\$2.5K
Growth (Mo 3-6)	100	1,200	\$1,400	\$500	\$900	\$8K
Steady (Mo 6-9)	250	3,000	\$3,500	\$1,100	\$2,400	\$22K
Scaling (Mo 9-12)	500	6,000	\$7,000	\$2,200	\$4,800	\$48K
Year 2 target	1,000	12,000	\$14,000	\$4,500	\$9,500	\$110K

Assumes Pro at \$14.
until cliff.

From break-even to founder salary, on a calendar

~15-20 Pro users	Break-even on cash costs	Week 4-6 if launch goes well
~50 Pro users	Founder pocket money — \$400-500/mo personal income	Month 2 (soft launch)
~250 Pro users	Replace one part-time consulting income (~\$2K/mo)	Month 6-9
~500 Pro users	Cover one founder's full salary (~\$5K/mo equiv.)	Month 9-12
~1,000 Pro users	Sustainable two-founder business with cash to reinvest	Year 2
~5,000 Pro users	Hire first FTE engineer; pursue B2B2C deals	Year 2-3

Per-user margin and fixed-cost coverage

- Per-Pro-user margin "H \$9-11 / month (after blended infra cost ~\$3-5).
- Fixed costs (domain, basic ops, eventually GitHub Team for cron capacity): ~\$50-200 / month until ~250 Pro users.
- Break-even: ~15-20 Pro users covers fixed costs.
- 100 Pro users = ~\$1,400 / month MRR · ~\$900 / month profit.
- 1,000 Pro users = ~\$14K / month MRR · ~\$9.5K / month profit.

Profit margin holds at 65-70% across scale.

When we leave each free tier and what it costs

Service	Free tier ceiling	First paid step	Trigger to upgrade
GitHub Actions	2,000 min / mo "H 250 Pro users	Self-hosted runner \$5–20/mo · GitHub Team \$4/seat	90% of free quota used 2 months in a row
Vercel	Hobby: 100 GB / mo "H 5,000 active users	Vercel Pro \$20/mo per seat	Bandwidth > 80 GB or function timeouts spiking
Firebase	Spark: 50K reads / 20K writes / day "H 600 active users	Blaze plan, scales with usage · ~\$50–200/mo at 5K users	Read or write quota hit 70% by mid-month
Apify	Free \$5/mo credit "H 10 Pro users	Personal \$49/mo (~200 users) · Team \$499/mo (~2,000)	Will hit immediately — budget from day 1
OpenRouter	Pay-as-you-go from \$0	Direct token cost — scales linearly	Cap if a user exceeds 10× median

Sequence: launch with Apify Personal
GitHub runner second (~250 Pro) !' Ve

What kills this if we are not paying attention

Risk 01. Apify quota or pricing changes

Mitigation: Build a second source (LinkedIn Jobs API, RemoteOK, WeWorkRemotely scraper) as cold-swap backup. Architecture already supports multiple discovery actors via the jobResearcher attempts loop.

Risk 02. GitHub free Actions ceiling at 2,000 min/mo

Mitigation: Monitor monthly usage; budget Team plan (\$4/seat) or self-hosted runner around the 200-user mark. Catch-up cron can drop primary to weekly if needed.

Risk 03. LLM cost volatility

Mitigation: Per-user monthly token budget enforced server-side. Auto-fallback to deterministic-only scoring if budget exceeded. Models swappable via OpenRouter without code change.

Risk 04. Match quality complaints

Mitigation: Weekly retro on dismiss / save / applied signals. AI scoring prompt is versioned. Per-company dedup tunable; prompt overrideable per career path family.

From launch to first 100 Pro users

Days 1 – 30 · Launch & instrument

- Ship Pro pricing at \$14/mo
- Wire cost telemetry per user (Apify + LLM tokens)
- 50 paying-user target
- Early-access list !' friends, ex-coworkers, comm groups

Days 31 – 60 · Build distribution

- Add second job source (LinkedIn / RemoteOK) for resilience
- Referral program: 1 free month per converted referral
- Product Hunt launch with 200 pre-registered upvoters
- 100 paying-user target

Days 61 – 90 · Scale & B2B

- Career coaches, bootcamps, layoff support orgs (B2B2C pilots)
- Annual plan promo (40% off launch window)
- Evaluate self-hosted GitHub runner
- 200 paying-user target

What we are deciding together

Capital

Initial runway for Apify + ad budget. Recommended: \$_____ for the first 90 days.

Time

Ongoing weekly commitment. Recommended: _____ hrs/week, focused on _____.

Distribution

Network access — career coaches, bootcamps, hiring communities, LinkedIn audience reach.

Sales

B2B2C lead conversion (career coaches / bootcamps). Recommended: pilot with ___ orgs by Day 60.

Structure to discuss: equity split · decision rights · vesting schedule · exit terms.

Thank you.

Three answers to leave with today:

1. Pricing

Lock \$14/mo Pro and \$99/yr annual?

2. Costs

Confirm the \$2.05–\$4.55 / Pro user / month basis is acceptable margin?

3. Break-even

Agree on 15-Pro-user break-even and 100-Pro-user 30-day target?